

# Project Proposal: Applying Graph Neural Networks in Molecular Property Prediction with PyTorch Geometric

Nikita Zagainov [n.zagainov@innopolis.university](mailto:n.zagainov@innopolis.university)  
Dmitry Tetkin [d.tetkin@innopolis.university](mailto:d.tetkin@innopolis.university)  
Nikita Tsukanov [n.tsukanov@innopolis.university](mailto:n.tsukanov@innopolis.university)

November 1, 2025

## Summary

We propose to develop a tutorial/case study on applying modern graph neural networks (GNNs) to molecular property regression in computational chemistry and drug discovery, using PyTorch Geometric (PyG). The tutorial will be step-by-step and self-contained for readers who know PyTorch but are new to graph ML, with clear code snippets, extensive visualizations, and Google Colab notebooks for full reproducibility.

## 1 Application domain

**Computational chemistry and drug discovery.** Molecules are naturally graphs: atoms as nodes and chemical bonds as edges. Predicting molecular properties is central to virtual screening and lead optimization.

## 2 Which dataset are you planning to use?

**Primary:** ZINC constrained solubility regression ( $\log P$ ), as curated by the GNN benchmarking suite (see [Benchmarking GNNs](#)); available directly in PyG as ZINC ([PyG ZINC](#)).

**Fallback/Optional extension:** QM9 quantum chemistry dataset ([PyG QM9](#); [Sci. Data 2014](#)).

## 3 Describe the dataset, prediction tasks, and metric

### Graph schema.

- **Nodes:** atoms, with categorical features such as element type, formal charge; optionally degree, aromaticity flags.
- **Edges:** chemical bonds, with bond type (single/double/triple/aromatic) as categorical edge features; optionally conjugation and ring membership.
- **Labels:** graph-level scalar property.

**Task:** supervised *graph-level regression*.

**Metric:** mean absolute error (MAE) on the target property.

**ZINC details:** *Constrained solubility* (logP) regression; small molecules (up to 28 atoms) curated for benchmarking ([Benchmarking GNNs](#)). Standard splits: 12k train / 1k val / 1k test with MAE as the official metric ([PyG ZINC](#); [paper](#)).

**QM9 details (optional):** Predict one or more quantum-mechanical properties (e.g., dipole moment, HOMO/LUMO, atomization energy). Standard MAE/% errors per-target ([PyG QM9](#); [dataset paper](#)).

## 4 Why did you choose the dataset?

- **Small graphs, fast training:** molecules are typically  $\downarrow$  30 nodes in ZINC, enabling rapid iteration even on modest GPUs.
- **Excellent PyG support:** convenient loaders and baselines reduce boilerplate ([PyG docs](#)).
- **Standardized splits and benchmarks:** ensures fair comparison and reproducibility ([Benchmarking GNNs](#); [benchmarking-gnns repo](#)).
- **Pedagogical clarity:** molecular graphs match the strengths of message passing, making concepts intuitive to newcomers.

## 5 Graph ML technique that you want to apply

**Message Passing Neural Networks (MPNNs)** with **edge-aware updates**. We will start with Graph Isomorphism Network (GIN; [paper](#)) and its edge-feature-aware variant GINE (available as GINEConv in PyG; [docs](#)). Pooling will use global add pooling, optionally with a virtual node.

## 6 Graph ML model you plan to use

**Backbone:** Stacked GINEConv layers with MLP message/update functions, each followed by BatchNorm, ReLU, and dropout.

**Readout:** Global add pooling to obtain a graph embedding, then a small MLP for scalar prediction.

**Training:** L2 loss (MSE), AdamW optimizer, cosine learning-rate schedule with warmup; early stopping on validation MSE.

**Baselines:** (i) RDKit descriptors + XGBoost ([RDKit](#); [XGBoost](#)), (ii) an MLP on simple atom-type histograms.

**Stretch (time-permitting):** virtual node, Stochastic Weight Averaging (SWA), and scaffold split robustness analysis.

## 7 Describe the model

**GIN/GINE (edge-aware MPNN).** We will use Graph Isomorphism Network (GIN; [paper](#)) and its edge-aware variant GINE ([PyG GINEConv](#)). Intuitively, each layer updates an atom by aggregating representations from its neighbors; GINE augments this with bond-type information so messages depend on both neighboring atoms and the connecting bond.

**Graph Attention (GAT/GATv2).** As a complementary model, we will evaluate attention-based layers that learn importance weights over neighbors: GAT ([paper](#), [PyG GATConv](#)) and GATv2 ([paper](#), [PyG GATv2Conv](#)). For molecules, edge features can modulate attention or messages, e.g., as in AttentiveFP ([paper](#)).

**Readout and training.** We use a permutation-invariant global add pooling to obtain a molecule-level embedding, followed by a small MLP for scalar prediction. Training minimizes mean squared error (MSE) with AdamW and cosine learning-rate scheduling; early stopping monitors validation MSE. For reporting and comparison to benchmarks, we will evaluate using mean absolute error (MAE), as specified in the dataset protocols.

**Planned architecture:** Input node/edge embeddings - $\downarrow$  GINEConv + BatchNorm + ReLU + Dropout (repeat) - $\downarrow$  Global Add Pooling - $\downarrow$  MLP Readout - $\downarrow$  Scalar prediction.

**Hyperparameters (to be tuned):** layers 3/5/7; hidden size 64/128/256; dropout 0.0–0.5; learning rate 1e-4–3e-3.

## 8 Why the model is appropriate for the dataset

- **Expressivity:** GIN matches the Weisfeiler–Lehman (1-WL) test in discriminative power ([Xu et al., 2019](#)), aligning with motif/substructure sensitivity needed for molecular properties.
- **Edge features:** GINE directly incorporates bond types and related chemistry, crucial for properties like solubility and electronic energies.
- **Data regime:** For small graphs and moderate dataset sizes like ZINC, lightweight MPNNs with global pooling are strong and efficient baselines.
- **Simplicity and reproducibility:** Popular, well-supported layers in PyG minimize engineering complexity while providing competitive performance.

## Links and references

- **PyTorch Geometric:** [docs](#)
- **ZINC in PyG:** [dataset](#)   **Benchmarking GNNs:** [paper](#)   [code](#)
- **QM9 in PyG:** [dataset](#)   **QM9 paper:** [Sci. Data 2014](#)
- **GIN (How Powerful Are GNNs?):** [paper](#)
- **GAT:** [paper](#)   [PyG GATConv](#)
- **GATv2:** [paper](#)   [PyG GATv2Conv](#)
- **AttentiveFP (edge-aware attention for molecules):** [paper](#)
- **GINEConv in PyG:** [docs](#)
- **Papers With Code (SOTA browser):** [link](#)   **OGB Leaderboards:** [link](#)
- **RDKit:** [link](#)   **XGBoost:** [link](#)

## Key references

- K. Xu, W. Hu, J. Leskovec, S. Jegelka, “How Powerful Are Graph Neural Networks?” ICLR 2019. [arXiv:1810.00826](https://arxiv.org/abs/1810.00826).
- P. Veličković et al., “Graph Attention Networks.” ICLR 2018. [arXiv:1710.10903](https://arxiv.org/abs/1710.10903).
- B. Brody, U. Alon, E. Yahav, “How Attentive are Graph Attention Networks?” ICLR 2022. [arXiv:2105.14491](https://arxiv.org/abs/2105.14491).
- S. Xiong et al., “Pushing the Boundaries of Molecular Representation for Drug Discovery with the Graph Attention Mechanism.” (AttentiveFP) [arXiv:1904.01279](https://arxiv.org/abs/1904.01279).
- V. Dwivedi et al., “Benchmarking Graph Neural Networks.” [arXiv:2003.00982](https://arxiv.org/abs/2003.00982).

*Group size:* 1–3 students. We will focus on well-established, classical graph ML methods and avoid research-heavy or novel architectures, prioritizing clarity and reproducibility.